

Table of Contents

Understanding What You Are Actually Buying

A plain-language guide to AI commercial constructs, with worked examples

Companion to the AI / ML Procurement Playbook, Part 1 | July 2026

Start here: the one idea

Every category you have ever sourced had a unit you could count before you signed. Licences. Seats. Devices. Hours. Litres.

AI does not have one.

You are buying **metered compute**. The meter counts a unit the vendor invented, the reading depends on engineering decisions made after you sign, and the same business outcome can cost ten times more or less depending on choices that appear nowhere in the contract.

That is the whole problem in one sentence. Everything below is a consequence of it.

There are five constructs. Vendors will blur them together in a single proposal. Your first job is to pull them apart.

#	Construct	The question it answers
1	Token pricing	How does the meter work?
2	Inference vs training	Which costs are ongoing and which are one-off?
3	Model licensing	What do we actually own or have a right to?
4	Data usage rights	What happens to our data, and who owns the output?
5	Hyperscaler AI services	Who do we buy it from, and on whose paper?

1. Token pricing

1.1 What a token is

A token is a fragment of text. Roughly **4 characters**, or about **three quarters of a word**.

“The vendor shall provide” is about 5 tokens. A single page of text is roughly 500 tokens. A 50-page contract is roughly 25,000 tokens.

Tokens are the billing unit. Prices are quoted per **million tokens**, written as MTok.

That is the entire concept. It is not complicated. What is complicated is everything the vendors do with it.

1.2 The two meters

Here is the first thing that trips up experienced buyers. There is not one meter. There are two, and they charge different rates.

Meter	What it counts	Relative price
Input	Everything you send: your question, the instructions, any documents attached, the conversation so far, definitions of any tools	Cheaper
Output	Everything the model writes back	5 to 6 times the input rate

Take a typical production model at \$3 per million input tokens and \$15 per million output tokens. Output costs five times what input costs.

Why this matters more than it sounds.

Imagine two use cases. Both process exactly 10,000 tokens per transaction. Both run 10,000 transactions a month. By any traditional procurement measure, they are the same size.

	Use case A: Summarize	Use case B: Draft
What it does	Reads a 9,000-token contract, writes a 1,000-token summary	Takes a 2,000-token instruction, writes an 8,000-token draft
Input tokens	9,000	2,000
Output tokens	1,000	8,000
Total tokens	10,000	10,000

Now the bill, at \$3 input and \$15 output, for 10,000 transactions:

	Use case A	Use case B
Input cost	90 MTok × \$3 = \$270	20 MTok × \$3 = \$60
Output cost	10 MTok × \$15 = \$150	80 MTok × \$15 = \$1,200
Monthly total	\$420	\$1,260

Identical volume. **Three times the cost.**

The buyer's rule: any vendor forecast that quotes you “tokens per month” as a single number is useless. Send it back and ask for the input and output split. If they cannot give you one, they have not sized your workload, they have guessed at it.

1.3 The sleeper cost: thinking tokens

Modern models do not answer immediately. They reason internally first, generating what the industry calls thinking or reasoning tokens, before producing the answer you actually see.

You are billed for that internal reasoning. At the output rate.

Worked example. A policy question chatbot:

	Tokens	Rate	Cost per 50,000 enquiries
Input (question + instructions + retrieved policy)	8,100	\$3 / MTok	\$1,215
Answer the user sees	400	\$15 / MTok	\$300
What you budgeted			\$1,515
Internal reasoning you	2,000	\$15 / MTok	\$1,500

	Tokens	Rate	Cost per 50,000 enquiries
never see			
What you actually pay			\$3,015

The reasoning is invisible, it is five times longer than the answer, and it doubles the bill. Nobody mentions it in the demo.

The question to ask every vendor: “Break your forecast into input tokens, output tokens, and reasoning tokens for our workload.” The silence that follows tells you how seriously they have sized it.

1.4 The four levers that move the bill

This is where procurement people usually stop reading and hand it to IT. Do not. These four levers are worth more than any discount you will ever negotiate, and understanding them is what lets you challenge a vendor’s forecast.

Lever 1: Caching (worth up to 90% of input cost)

Most AI applications send the same thing over and over. Our policy chatbot sends the same 2,000-token instruction set and the same 6,000 tokens of retrieved policy on every single enquiry. Only the 100-token question changes.

Caching means: the vendor stores that repeated chunk and charges roughly **10% of the normal input rate** to read it back, instead of reprocessing it from scratch every time.

	Without caching	With caching
Repeated portion (8,000 tokens × 50,000)	400 MTok × \$3 = \$1,200	400 MTok × \$0.30 = \$120
The actual question (100 tokens × 50,000)	5 MTok × \$3 = \$15	5 MTok × \$3 = \$15
Input cost	\$1,215	\$135

That is 89% off the input side of the bill, for free. It costs a small amount to write to the cache in the first place, which is immaterial at this volume.

This is not a negotiation. It is an engineering choice. No vendor will offer it to you. You have to require it.

Lever 2: Batching (worth 50%)

If your work can wait up to 24 hours for an answer, every major vendor will do it for **half price**. Flat 50% off both meters.

Nightly reporting, bulk document processing, historical analysis, evaluation runs, back-office summarization. None of this needs an answer in two seconds.

Batching stacks with caching. A cached and batched workload can run at a small fraction of list price.

Lever 3: Model tier (worth 3x to 25x)

This is the big one. Vendors sell a ladder:

Tier	Example rate	Good for
Frontier	\$10 / \$50	The genuinely hardest reasoning
Flagship	\$5 / \$25	Complex analysis, difficult judgement
Production	\$2 / \$10 to \$3 / \$15	The overwhelming majority of real work
High-volume	\$1 / \$5	Classification, routing, extraction, simple answers
Ultra-budget	\$0.20 / \$1.25	Sorting, tagging, triage

The gap between the top and bottom of that ladder is **25 to 50 times**. The gap between two vendors’ flagships is about 20%.

Ask yourself which of those two numbers you should be spending your negotiating energy on.

Most enterprise workloads do not need a flagship model. “Which of these 200 invoices mention a penalty clause” is not a hard reasoning problem. Route it to the cheap tier. Reserve the expensive tier for the 10% of tasks that genuinely need it.

Lever 4: Endpoint and residency (costs you about 10%)

If you require inference to stay in a named geography, that typically costs about a **10% uplift**. For Ontario public sector work with data residency requirements, this is not optional and it belongs in your base case from day one, not discovered in month three.

The four levers stacked

Our 50,000-enquiry chatbot, on the production tier, before and after:

Configuration	Monthly cost
Flagship model, no caching, no batching, all real-time	~\$5,000
Production tier, no caching	~\$3,000
Production tier, caching enabled	~\$1,900
Production tier, caching, 50% of volume batched	~\$1,300
Simple enquiries routed to high-volume tier, caching, batching	~\$500

Figures illustrative and rounded, based on the token shapes above.

Ten times the difference. Zero negotiation. All architecture.

The single most important sentence in this document: do the architecture work before the commercial work. If you go to the negotiating table before your team has done model routing and caching, you are fighting hard for 20% while leaving 80% on the table.

2. Inference vs training

2.1 The distinction, simply

	Inference	Training / fine-tuning
Plain meaning	Asking the model a question and getting an answer	Teaching the model something about your organization
When it happens	Every single time anyone uses it, forever	Once, or occasionally
Cost shape	Variable, ongoing, grows with adoption	Looks like a one-time project cost
Feels like	A utility bill	A capital project

They belong in different lines of the business case. They almost never are.

2.2 The fine-tuning trap

Here is a conversation that happens in every organization.

Vendor: “We can fine-tune a model on your ten years of procurement documents. It will understand your terminology, your templates, your directives.”

You: “What does that cost?”

Vendor: “The training run is a few hundred dollars.”

A few hundred dollars. For an AI that knows your business. Who would say no?

Here is what was not said.

A fine-tuned model is *your* model. It is not shared with anyone else. Which means it needs *its own capacity to sit on*. And that capacity is billed **by the hour, whether or not anybody uses it**.

The training run	A few hundred dollars, one time
Hosting the model, 24 hours a day, 365 days a year	Ask. This is the real number.
What happens when the base model is deprecated in 9 to 12 months	You retrain. At your cost.
Effective cost over 3 years	An annuity, not a project

The “few hundred dollar” fine-tune is a multi-year recurring obligation wearing a one-time cost’s clothing.

Ask this, in writing: “What is the total cost to keep this fine-tuned model deployed and reachable for 36 months, and what happens to it when the base model is retired?”

2.3 The alternative nobody offers you

There is a technique called **RAG**, which stands for retrieval-augmented generation. In plain language: instead of teaching the model your documents permanently, you just *hand it the relevant documents at the moment you ask the question*.

The comparison:

	Fine-tuning	RAG
How it works	Retrain the model on your data	Look up the relevant documents, include them in the question
Upfront cost	Training run	Building a retrieval system
Ongoing cost	Dedicated hosting, forever	Just input tokens, which are cacheable
Updating it	Retrain	Update a document. Done in seconds
When the base model changes	Retrain from scratch	Change one setting
Lock-in	High. The model is vendor-specific	Low. Your documents are yours
Auditability	Hard to explain why it said something	You can see exactly which document it used

For the overwhelming majority of enterprise use cases, **RAG is better, cheaper, more auditable, and more current**. Fine-tuning is right when you need the model to adopt a *style* or a *format* it cannot be instructed into, which is rarer than vendors suggest.

Notice the auditability row. For an Ontario public body that may have to explain a decision to an affected person or a regulator, “the model just knows” is not an answer. “It used these three documents” is.

The counter to make: “Justify fine-tuning against a RAG baseline. Show me the evaluation results for both.”

If a vendor proposes fine-tuning early and cannot answer that, they may be selling you lock-in and a hosting annuity.

3. Model licensing

3.1 The three archetypes

Think of it as three ways to get from A to B.

	Proprietary API	Open-weight	Embedded in SaaS
Analogy	Taking a taxi. You pay per trip. You never own the car. You cannot look under the bonnet	Buying a car. The car is cheap or free. Fuel, garage, insurance, and a mechanic are on you	A hotel that added free breakfast and raised the room rate. You cannot see what the eggs cost
Examples	Claude, GPT, Gemini	Llama, Mistral, DeepSeek, Qwen, Gemma	Copilot, an AI feature in your CRM or EMR
What you get	A right to call an endpoint	The model file itself	A feature
You pay	Per token	For your own compute and staff	Per seat, or per credit, or it is bundled
Meter visibility	Full	Total, it is your infrastructure	None
Lock-in	Moderate	Low on the vendor, high on your own build	Highest

3.2 On “open” models

“Open” is a marketing word, not a legal one. Treat an open-weight licence exactly as you would treat an open-source software licence review, because that is what it is.

Things that routinely catch buyers out:

- A **user threshold** above which you must stop and negotiate a commercial licence
- An **acceptable use policy** incorporated by reference, which the licensor can change unilaterally
- A ban on using the outputs **to train a competing model**
- **Attribution and naming** obligations
- **Field-of-use** restrictions

And the commercial reality: open-weight is only cheap at very high, very steady volume, or when data legally cannot leave your estate. At low volume you are paying for idle GPUs and an MLOps team you probably do not have. “Free model” and “cheap solution” are entirely different claims.

3.3 Embedded AI: the one that will actually get you

This is where most public sector AI spend will land, and it is the hardest to govern.

Here is how it happens. Your CRM vendor, or your EMR vendor, or your HR system vendor sends a routine terms update. Buried in it: AI features are now enabled.

In that moment, with no procurement process, no evaluation, and no business case, your organization has acquired an AI system. You have inherited:

- Their choice of model, which they can change without telling you
- Their subprocessors, in their jurisdictions
- Their data flows
- Their pricing, on a meter you cannot read
- An AI use case that, under the OPS AI Directive, you may now be obliged to assess, register, and disclose

And your leverage is at renewal, and nowhere else.

Practical action: ask your major application vendors, today, in writing: “Have you enabled AI features in our instance? Which model, hosted where, and does our data train it?” The answers will be more interesting than you expect.

4. Data usage rights

This is the clause set that separates a competent AI buyer from an incompetent one.

There are **six** distinct rights. Vendors typically address two of them clearly and hope you do not ask about the other four.

4.1 The six rights

#	The right	What goes wrong if you skip it
1	No training on our inputs	Your prompts and documents improve a model your competitors and the public then use
2	No training on our outputs	Same problem, different door
3	Human review narrowed	A vendor's contractor reads a citizen's health information during abuse monitoring
4	Retention and residency defined	Your "Canadian" data sits on a foreign server for 30 days under an abuse-monitoring policy nobody read
5	Output ownership	Ambiguity about who owns work product your organization is publishing
6	IP indemnity that survives contact with reality	A third party claims the output infringed their copyright and you are alone

4.2 Right 1 and 2: training on your data

What vendors say: "We do not train on enterprise customer data."

What that is: a policy, published on a webpage, that they can change.

What you need: a contract term, covering inputs, outputs, and derived data, extending to every subprocessor, with an audit right.

The distance between those two is the entire job.

The reply that works: "Excellent. Since that is already your policy, you will have no difficulty putting it in the agreement."

The response to *that* sentence tells you everything.

4.3 Right 3: human review

Almost every vendor reserves the right for humans to look at your data. The stated reason is usually abuse detection or safety, and it is a legitimate one.

The problem is the drafting. A clause permitting review for "service improvement" is a clause permitting review for anything.

The realistic landing zone:

	Position
Trigger	Only on a documented security or abuse signal
Scope	Only the flagged content, not the account
Notice	Where legally permitted
Logging	Every access logged and available to you

	Position
Personnel	Named categories, background checked, bound by confidentiality

4.4 Right 4: retention and residency

Watch this sequence:

1. The vendor confirms data residency in Canada. Your PIA is approved.
2. Their standard terms include 30-day retention “for abuse monitoring.”
3. That retained copy sits in the vendor’s central systems, not the Canadian region.
4. Your residency guarantee applied to *inference*. It did not apply to *retention*.

Nobody lied. You just did not ask the second question.

Ask all four:

- Where is the data processed?
- Where is it **retained**, and for how long?
- Who else touches it, and where are they?
- Can we get certified deletion, on demand and on termination?

4.5 Right 5: output ownership

Most vendors will grant you ownership of outputs. Good. Read the sentence that follows.

Ownership is frequently made **conditional** on your compliance with an acceptable use policy that is (a) on a webpage and (b) unilaterally amendable.

That means the vendor can, in principle, change the conditions of your ownership after you have signed.

Get ownership clean and unconditional. Accept AUP compliance as a separate covenant with its own remedy. Do not let it be a condition of title.

4.6 Right 6: the indemnity, and why the cap is not the point

Every procurement person instinctively negotiates the **cap**. On AI deals, the cap is the wrong fight.

Here is a typical AI output indemnity, in plain language:

*“We will defend you against third-party intellectual property claims arising from outputs, **provided that you used our default safety settings and did not modify the output.**”*

Now think about your actual workflow. Somebody takes the AI’s draft, edits it, and puts it in a document.

You just modified the output. The indemnity is void.

It does not matter whether the cap was \$1 million or \$50 million. You voided it through completely normal use, on day one, and nobody in the organization noticed.

The lesson: on AI contracts, **attack the conditions before you argue about the cap**. An indemnity you void by using the product as intended is not an indemnity. It is decoration.

4.7 The Ontario overlay

Two things frame everything above for your context:

Instrument	What it is
Responsible Use of Artificial Intelligence Directive	The OPS directive setting requirements for transparent,

Instrument	What it is
(effective 1 December 2024)	responsible, accountable AI use. Covers procured, internally built, and publicly available tools, across the full lifecycle
IPC-OHRC Principles for the Responsible Use of AI (January 2026)	Joint guidance from the privacy and human rights regulators. Not law, but explicitly the standard those regulators say they will use to assess you

The IPC-OHRC six principles are: valid and reliable, safe, privacy-protective, human rights affirming, transparent, and accountable.

Use them as your evaluation criteria structure. It is defensible, it maps to what a regulator will actually ask, and it saves you inventing your own framework.

And the practical point that matters most: under FIPPA and PHIPA you need lawful authority and real controls before personal or health information goes to a processor. “The vendor says they do not train on it” is not a control. A contract term with an audit right and certified deletion is a control.

5. Hyperscaler AI services

5.1 What you are actually buying

The same models are available directly from the model vendor, and through AWS, Microsoft, and Google.

Here is the thing nobody says out loud: **you are often not buying cheaper tokens**. On several platforms the rate is the same as buying direct.

So what are you paying for?

What the hyperscaler route gives you	Why it is worth real money
Existing paper	No new MSA, no new legal review, no new vendor onboarding
Existing security posture	Your cloud security review already covers the tenant
Existing procurement vehicle	You may already have a route to market
Drawdown against an existing commitment	The spend consumes money you have already committed to spend
Governance stack	Your existing monitoring, policy, and cost management tools apply
Residency controls	Regional endpoints and data zone options

That fourth row is the one that matters. If your organization has already committed to a cloud spend, routing AI through it means the AI spend is coming out of a pot you were going to spend anyway. Benchmarking work suggests that routing through an existing commitment unlocks **10% to 25% more discount** than a standalone deal.

Be honest about this internally. You are buying a commercial wrapper, not a cheaper token. That is a perfectly good reason to buy. Just do not let a vendor sell it to you as a discount, and do not let your own business case claim a saving that is not there.

5.2 The three routes, briefly

	AWS Bedrock	Microsoft Azure / Foundry	Google Vertex
Character	A marketplace. Many model vendors	OpenAI models plus a growing catalogue, inside	Google’s own models plus others, on Google’s own

	AWS Bedrock	Microsoft Azure / Foundry	Google Vertex
Buy it if	You are on AWS and want model optionality	your Azure tenant You are on an EA and want the spend to draw down a commitment	stack You are on GCP and have long-context or multimodal workloads
Commercial hook	Marketplace private offers, Savings Plans on underlying compute	EA / MCA, counts toward your Microsoft commitment	Committed use discounts

5.3 Reserved capacity: where AI budgets go to die

Every hyperscaler sells reserved capacity. Microsoft calls the unit a **PTU** (Provisioned Throughput Unit). AWS and Google have equivalents.

The pitch: commit to dedicated capacity, get a lower effective rate per token and guaranteed, consistent performance.

The catch, and it is a big one: *you pay for the capacity whether or not you use it.* It is not per token any more. It is a monthly floor. You cannot pause it. Billing stops when you delete the deployment, not when you stop using it.

The worked example:

You commit to reserved capacity costing	\$10,000 / month
Your actual utilization turns out to be	30%
Value of work actually done	\$3,000 / month
What consumption pricing would have cost	\$3,000 / month
Waste	\$7,000 / month, or \$84,000 / year

That 30% is not a pessimistic invention. Benchmarking across enterprise engagements consistently finds that reserved throughput bought **before** a workload has real production traffic data runs at **20% to 40% utilization**.

The published break-even guidance is roughly: reserved capacity wins when you are sustaining very high throughput (on the order of 300,000 tokens per minute) for **8 or more hours a day**, consistently. Below that, consumption pricing wins.

The correct sequence, which is the opposite of what you will be told:

Step	Action
1	Run on consumption pricing . No commitment.
2	Instrument everything . Measure for 60 to 90 days.
3	Look at your P95 sustained throughput , not your peak, and not your forecast.
4	Convert only the proven steady-state load to reserved capacity.
5	Do it at renewal , when the commitment is also a negotiation asset.
6	Leave the peak on consumption.

The counter when a vendor pushes reserved capacity into a pilot: “Show me the telemetry that justifies it. We will convert the proven steady state at renewal.”

A vendor pushing reserved capacity before you have production traffic is selling you idle capacity.

5.4 The costs that are not on the price list

Teams that budget only for tokens routinely land **15% to 40% over plan**. Here is where it goes:

Hidden cost	What it is
Data egress	Moving data between clouds. A retrieval pipeline pulling documents from storage in one cloud to a model in another bleeds here, per gigabyte, forever
Vector database	The retrieval infrastructure for RAG. Its own bill
Guardrails and content filtering	Often a separate meter
Monitoring and observability	You cannot manage what you cannot see, and seeing costs money
Private networking	If you need it, it is not free
Support tier uplift	The support level you actually need is rarely the included one
Idle fine-tuned models	See section 2.2. The zombie that bills forever
Human evaluation and review	Real FTE cost. Never in a vendor quote. Always in your reality

Put a named contingency in the business case and defend it with this list.

6. Ten things to carry into the room

#		
1	Output costs 5x input.	Two workloads of identical size can differ 3x in cost. Always get the split.
2	Thinking tokens are billed as output.	They are invisible, often longer than the answer, and can double your bill.
3	Architecture beats negotiation.	Model tier, caching, and batching are worth 3x to 10x. Discounts are worth 15% to 30%.
4	Cost per token is the wrong metric.	Cost per successful outcome is the right one. A model costing 20% more that needs 40% fewer retries is cheaper.
5	Fine-tuning is a hosting annuity.	The training run is not the cost. Ask for 36 months of hosting, and what happens at depreciation.
6	Try RAG first.	Cheaper, more current, more auditable, less lock-in. Make the vendor justify fine-tuning against it.
7	“Open” is a marketing word.	Read the licence like an OSS licence. Check the user thresholds.
8	A policy is not a contract term.	“We do not train on your data” belongs in the agreement, not on a webpage.
9	Attack indemnity conditions, not the cap.	An indemnity you void by editing an output is decoration.
10	Reserved capacity before telemetry is waste.	Consumption first. Measure 60 to 90 days. Convert the proven steady state at renewal.

Appendix: the ten questions

Put these to any AI vendor, in writing, and read the answers carefully.

1. Split your forecast for our workload into input tokens, output tokens, and reasoning tokens.
 2. What percentage of our input do you expect to be cacheable, and on what basis?
 3. What is the cost per successful outcome, on our benchmark set, at each model tier you offer?
 4. Which of your models have you deprecated in the last 24 months, and what notice did customers get?
 5. Does the model you are quoting use the same tokenizer as the model we benchmarked?
 6. What are we billed if we send a request and receive an error?
 7. What is your uptime commitment, what is the remedy, and how many times have you missed it in the last 12 months?
 8. Name every subprocessor that will touch our data, and the country each operates from.
 9. Under exactly what circumstances will a human at your organization read our data?
 10. What does it cost us to leave, and what do we get back?
-

Pricing illustrations use representative market rates as at July 2026 and are for teaching purposes. Verify current rates against vendor documentation before any figure enters a business case. This is a commercial and procurement guide, not legal advice.